Large-scale data analysis in R

Research Computing Center and the Dept. of Human Genetics University of Chicago

Peter Carbonetto
pcarbo@uchicago.edu

November 14, 2017

Getting started

Initial setup:

- Wifi
- Power outlets
- RCC cluster access
- Yubikeys (`rcc-guest-xxxx`)
- Helper(s)
- Handout / exercises
- Pace, questions (e.g., keyboard shortcuts)

Getting started

1. Download workshop packet.

```
$ cd ~
$ wget https://tinyurl.com/yad63z9c \
    -O temp.tar.gz
$ tar zxvf temp.tar.gz
$ cd R-large-scale
$ cd code
```

2. Add these three lines to your `.bashrc` file.

```
# SLURM.
export SACCT_FORMAT="jobid,partition,user,account%12,alloccpus,node%12,elapsed,totalcpu,maxRSS,ReqM"
export SQUEUE_FORMAT="%13i %12j %10P %10u %12a %8T %9r %10l %.11L %5D %4C %8m %N"
```

Getting started

3. Start up three midway2 sessions.

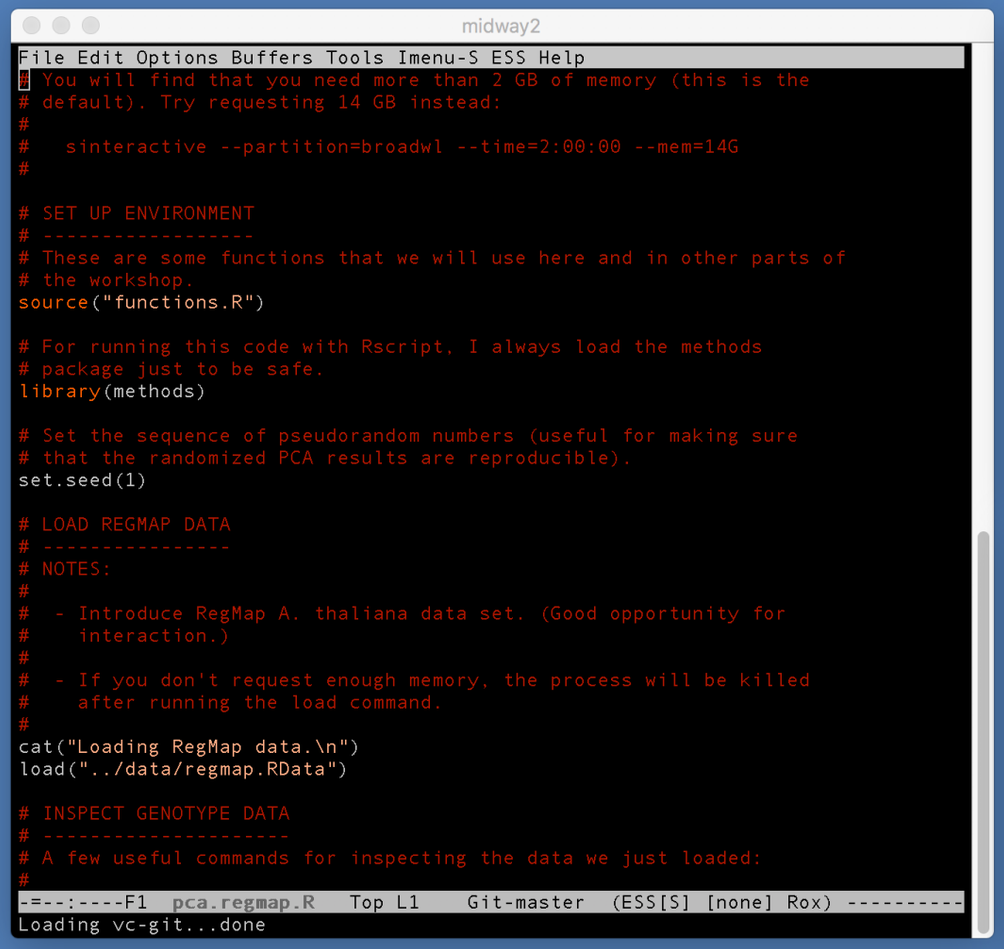

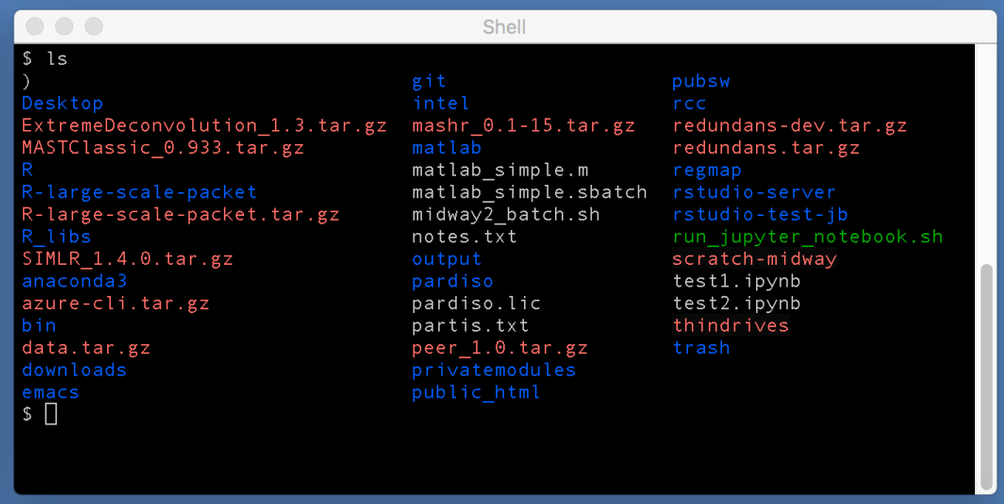

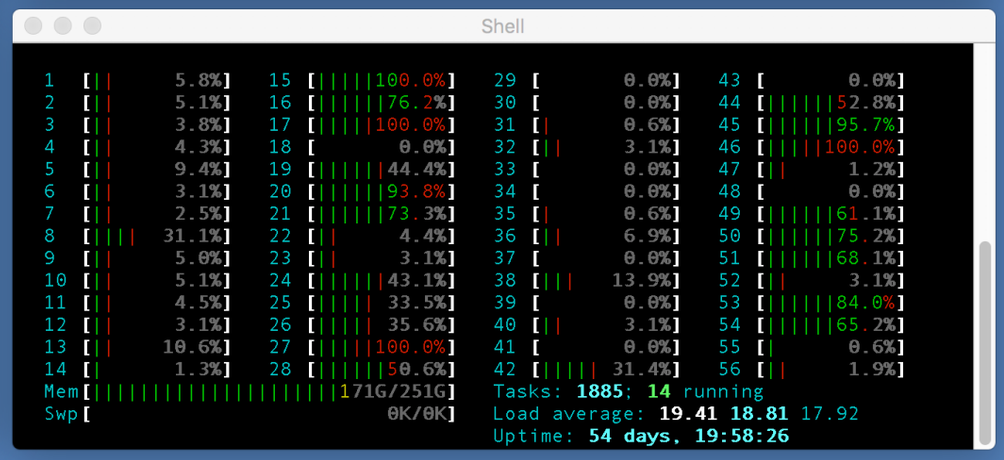

Getting started

4. Request a compute node.

```
$ screen -S rcc_workshop
$ sinteractive --partition=broadwl \
    --time=2:00:00 --account=rcc-guest \
    --reservation=rworkshop2
$ echo $HOSTNAME
```

5. Start up R programming environment.

```
$ module load R/3.4.1
$ R
R> getwd()
```

Aims of workshop

- ~~Do an analysis of a large-scale data set in R.~~
- Develop some useful skills for large-scale data analysis in R *within a high-performance computing environment*, and apply these skills to a medium-scale data set.

Outline

- Initial setup.
- Brief introduction.
- Part 1: Warm-up with PCA of RegMap data.
- Part 2: Implementing multithreaded computation in R for analysis of genetic adaptation to climate.
- Brief recap.

Part 1: PCA of the RegMap data

Adaptation to Climate Across the *Arabidopsis thaliana* Genome

Angela M. Hancock,[1] Benjamin Brachi,[2] Nathalie Faure,[2] Matthew W. Horton,[1]
Lucien B. Jarymowycz,[1] F. Gianluca Sperone,[1] Chris Toomajian,[3] Fabrice Roux,[2] Joy Bergelson[1]*

www.sciencemag.org **SCIENCE** VOL 334 7 OCTOBER 2011

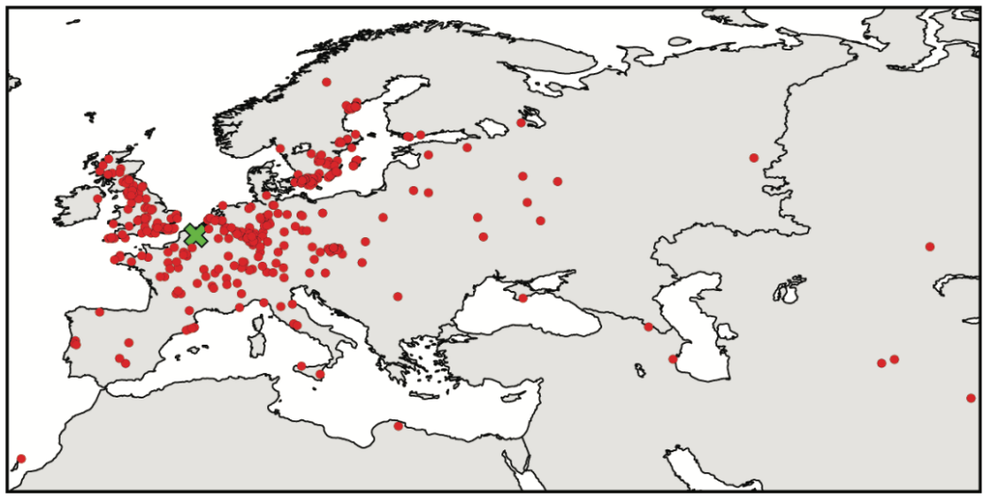

1. Inspect RegMap data.
2. Run PCA interactively and assess time & memory needs.
3. Examine PCA results.
4. Run PCA using Rscript.
5. Run PCA using SLURM job engine, and inspect output.

Exercise: Run PCA using SLURM job engine, and inspect output.

```
sbatch pca.sbatch
```

- Useful commands while job is running:

```
squeue --user=<cneitd> | less -S
ssh midway2-xxxx
```

- Useful commands after job has completed:

```
htop --user=<cnetid>
sacct --user=<cnetid> --units=G | less -S
less ../output/pca_err.txt
less ../output/pca_out.txt
```

Part 2: Genetic analysis of climate variables in RegMap data

1. Inspect climate ("phenotype") data.
2. Run analysis interactively and assess time & memory needs.
3. Experiment with multithreaded matrix operations to improve computation time.
4. Experiment with parallel computation of the weights using `mclapply`.
5. Automate `climate.R` analysis using command-line arguments.
6. Run analysis using SLURM job engine, and inspect output.
7. Use SLURM job engine to automate analysis of all 48 climate variables.

There is no best tool—use *whatever works for you*.

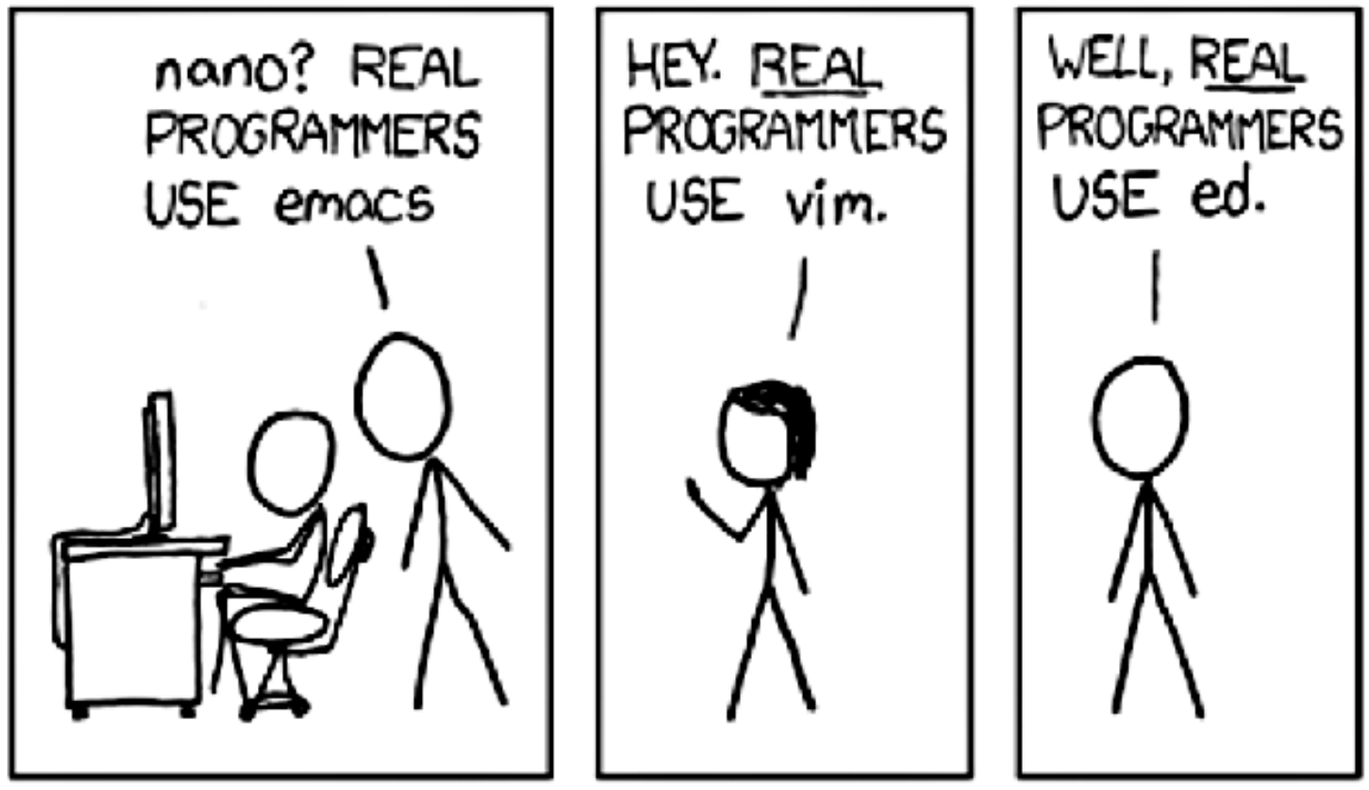

https://xkcd.com/378/

Some general advice

1. `help(package = cool_package)`.
2. Use **midway2**, not midway1.
3. Email help@rcc.uchicago.edu R help on the RCC cluster.
4. Learn to avoid loops as much as possible; e.g., use `apply()`, `lapply()`, `tapply()`, `do.call()`.
5. Document your setup—start with `sessionInfo()`.

After the workshop

- Feedback.
- Suggestions for future R workshop topics (e.g., Rcpp).